

# PROJECT REPORT

## *CLASS PROJECT PART – 1*

### INTRODUCTION AND PROBLEM DESCRIPTION

This project focuses on using the analytical skills by using big data technologies like AWS S3, AWS EMR – Hive and HDFS and AWS Athena. We will be using the Amazon reviews dataset available in S3. Our dataset will be in parquet format to improve our processing and is partitioned by product category. We begin our analysis from 2005. We choose a few categories from the list available and exclude multiple reviews by customers and only choose the most recent reviews. Based on the results, we will attempt to answer simple data exploratory questions along with trend analysis of our metrics. We will also attempt to correlate multiple product categories and answer questions based on the results.

We aim to perform analysis on the amazon reviews dataset. We will attempt to analyse trends on these metrics and try to correlate them and use different window functions and analytical aggregate functions to demonstrate the concepts of percentiles and moving average. We will also visualize some of our findings to provide clarity on the results obtained.

We begin by provisioning the EMR cluster and then copying the amazon reviews to EMR's HDFS. We then create an external table in Hive by pointing it to the HDFS folder. We then perform query operation on this table.

### TECHNICAL SCRIPTS, SQL QUERIES AND EXPLANATION WITH VISUALIZATIONS

The steps we will perform to run hive queries on the external table are as follows:

1. Create directories in HDFS for our product categories

Query:

```
`hdfs dfs -mkdir -p /hive/amazon-reviews-pds/parquet/product_category=Wireless/`  
`hdfs dfs -mkdir -p /hive/amazon-reviews-pds/parquet/product_category=Automotive/`  
`hdfs dfs -mkdir -p /hive/amazon-reviews-pds/parquet/product_category=Music/`  
`hdfs dfs -mkdir -p /hive/amazon-reviews-pds/parquet/product_category=Digital_Music_Purchase/`  
`hdfs dfs -mkdir -p /hive/amazon-reviews-pds/parquet/product_category=Sports/`  
`hdfs dfs -mkdir -p /hive/amazon-reviews-pds/parquet/product_category=Toys/`  
`hdfs dfs -mkdir -p /hive/amazon-reviews-pds/parquet/product_category=Digital_Video_Games/`  
`hdfs dfs -mkdir -p /hive/amazon-reviews-pds/parquet/product_category=Video_Games/`
```

```
[hadoop@ip-172-31-79-160 ~]$ hdfs dfs -mkdir -p /hive/amazon-reviews-pds/parquet/product_category=Wireless/
hdfs dfs -mkdir -p /hive/amazon-reviews-pds/parquet/product_category=Automotive/
hdfs dfs -mkdir -p /hive/amazon-reviews-pds/parquet/product_category=Music/
hdfs dfs -mkdir -p /hive/amazon-reviews-pds/parquet/product_category=Digital_Music_Purchase/
hdfs dfs -mkdir -p /hive/amazon-reviews-pds/parquet/product_category=Sports/
hdfs dfs -mkdir -p /hive/amazon-reviews-pds/parquet/product_category=Toys/
hdfs dfs -mkdir -p /hive/amazon-reviews-pds/parquet/product_category=Digital_Video_Games/
hdfs dfs -mkdir -p /hive/amazon-reviews-pds/parquet/product_category=Video_Games/
[hadoop@ip-172-31-79-160 ~]$ hdfs dfs -mkdir -p /hive/amazon-reviews-pds/parquet/product_category=Automotive/
[hadoop@ip-172-31-79-160 ~]$ hdfs dfs -mkdir -p /hive/amazon-reviews-pds/parquet/product_category=Music/
[hadoop@ip-172-31-79-160 ~]$ hdfs dfs -mkdir -p /hive/amazon-reviews-pds/parquet/product_category=Digital_Music_Purchase/
[hadoop@ip-172-31-79-160 ~]$ hdfs dfs -mkdir -p /hive/amazon-reviews-pds/parquet/product_category=Sports/
[hadoop@ip-172-31-79-160 ~]$ hdfs dfs -mkdir -p /hive/amazon-reviews-pds/parquet/product_category=Toys/
[hadoop@ip-172-31-79-160 ~]$ hdfs dfs -mkdir -p /hive/amazon-reviews-pds/parquet/product_category=Digital_Video_Games/
[hadoop@ip-172-31-79-160 ~]$ hdfs dfs -mkdir -p /hive/amazon-reviews-pds/parquet/product_category=Video_Games/
[hadoop@ip-172-31-79-160 ~]$
```

- Copy data from S3 to these directories

Query:

```
`s3-dist-cp --src=s3://amazon-reviews-pds/parquet/product_category=Wireless/ --
dest=hdfs:///hive/amazon-reviews-pds/parquet/product_category=Wireless/`
`s3-dist-cp --src=s3://amazon-reviews-pds/parquet/product_category=Automotive/ --
dest=hdfs:///hive/amazon-reviews-pds/parquet/product_category=Automotive/`
`s3-dist-cp --src=s3://amazon-reviews-pds/parquet/product_category=Music/ --
dest=hdfs:///hive/amazon-reviews-pds/parquet/product_category=Music/`
`s3-dist-cp --src=s3://amazon-reviews-pds/parquet/product_category=Digital_Music_Purchase/ --
dest=hdfs:///hive/amazon-reviews-pds/parquet/product_category=Digital_Music_Purchase/`
`s3-dist-cp --src=s3://amazon-reviews-pds/parquet/product_category=Sports/ --
dest=hdfs:///hive/amazon-reviews-pds/parquet/product_category=Sports/`
`s3-dist-cp --src=s3://amazon-reviews-pds/parquet/product_category=Toys/ --
dest=hdfs:///hive/amazon-reviews-pds/parquet/product_category=Toys/`
`s3-dist-cp --src=s3://amazon-reviews-pds/parquet/product_category=Digital_Video_Games/ --
dest=hdfs:///hive/amazon-reviews-pds/parquet/product_category=Digital_Video_Games/`
`s3-dist-cp --src=s3://amazon-reviews-pds/parquet/product_category=Video_Games/ --
dest=hdfs:///hive/amazon-reviews-pds/parquet/product_category=Video_Games/`
```

```
Map output materialized bytes=692
Input split bytes=155
Combine input records=0
Combine output records=0
Reduce input groups=10
Reduce shuffle bytes=692
Reduce input records=10
Reduce output records=0
Spilled Records=20
Shuffled Maps =3
Failed Shuffles=0
Merged Map outputs=3
GC time elapsed (ms)=842
CPU time spent (ms)=25780
Physical memory (bytes) snapshot=1544613888
Virtual memory (bytes) snapshot=17144954880
Total committed heap usage (bytes)=1383596032

Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters
Bytes Read=3326
File Output Format Counters
Bytes Written=0
/04/13 01:32:20 INFO s3distcp.S3DistCp: Try to recursively delete hdfs:/tmp/70
3133-a97f-4e2e-a673-250e2a734ac4
[hadoop@ip-172-31-79-160 ~]$ s3-dist-cp --src=s3://amazon-reviews-pds/parquet/pr
duct_category=Sports/ --dest=hdfs:///hive/amazon-reviews-pds/parquet/product_ca
tegory=Sports/
/04/13 01:32:21 INFO s3distcp.S3DistCp: Running with args: -libjars /usr/share
aws/emr/s3-dist-cp/lib/byte-buddy-1.9.10.jar,/usr/share/aws/emr/s3-dist-cp/lib/
te-buddy-agent-1.9.10.jar,/usr/share/aws/emr/s3-dist-cp/lib/commons-httpclient
.1.jar,/usr/share/aws/emr/s3-dist-cp/lib/commons-logging-1.0.4.jar,/usr/share/
s/emr/s3-dist-cp/lib/guava-18.0.jar,/usr/share/aws/emr/s3-dist-cp/lib/mockito-
re-2.27.0.jar,/usr/share/aws/emr/s3-dist-cp/lib/objenesis-2.6.jar,/usr/share/a
```

## 3. Create database

Query:

```
`create database amazon_review;`
```

```
[hadoop@ip-172-31-79-160 ~]$ hive
Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j2.properties Async: false
hive> create database amazon_review;
OK
Time taken: 0.772 seconds
hive> use amazon_review;
OK
Time taken: 0.04 seconds
```

## 4. Create an external table pointing to HDFS

Query:

```
`CREATE EXTERNAL TABLE amazon_review.amazon_reviews_parquet(
  `marketplace` string,
  `customer_id` string,
  `review_id` string,
  `product_id` string,
  `product_parent` string,
  `product_title` string,
  `star_rating` int,
  `helpful_votes` int,
  `total_votes` int,
  `vine` string,
  `verified_purchase` string,
  `review_headline` string,
  `review_body` string,
  `review_date` DATE,
  `year` int)
PARTITIONED BY (
  `product_category` string)
--ROW FORMAT DELIMITED
--STORED AS PARQUET
ROW FORMAT SERDE
'org.apache.hadoop.hive.ql.io.parquet.serde.ParquetHiveSerDe'
STORED AS INPUTFORMAT
'org.apache.hadoop.hive.ql.io.parquet.MapredParquetInputFormat'
OUTPUTFORMAT
'org.apache.hadoop.hive.ql.io.parquet.MapredParquetOutputFormat'
LOCATION
'hdfs:///hive/amazon-reviews-pds/parquet/'
TBLPROPERTIES (
  'transient_lastDdlTime'='1583454851');`
```

```
Msck repair table amazon_review.amazon_reviews_parquet;
```

```

Time taken: 0.502 seconds
hive> CREATE EXTERNAL TABLE amazon_review.amazon_reviews_parquet(
> `marketplace` string,
> `customer_id` string,
> `review_id` string,
> `product_id` string,
> `product_parent` string,
> `product_title` string,
> `star_rating` int,
> `helpful_votes` int,
> `total_votes` int,
> `vine` string,
> `verified_purchase` string,
> `review_headline` string,
> `review_body` string,
> `review_date` DATE,
> `year` int)
> PARTITIONED BY (
> `product_category` string)
> --ROW FORMAT DELIMITED
> --STORED AS PARQUET
> ROW FORMAT SERDE
> 'org.apache.hadoop.hive.ql.io.parquet.serde.ParquetHiveSerDe'
> STORED AS INPUTFORMAT
> 'org.apache.hadoop.hive.ql.io.parquet.MapredParquetInputFormat'
> OUTPUTFORMAT
> 'org.apache.hadoop.hive.ql.io.parquet.MapredParquetOutputFormat'
> LOCATION
> 'hdfs:///hive/amazon-reviews-pds/parquet/'
> TBLPROPERTIES (
> 'transient_lastDdlTime'='1583454851');
OK

```

5. Create an exclude table to exclude multiple reviews by a customer

We create a table that will exclude all the review ids for customers that have multiple reviews.

Query:

```

`CREATE TABLE amazon_reviews_exclude AS`
`select marketplace, customer_id, review_id, product_id, product_parent, product_title, star_rating,
helpful_votes, total_votes, vine, verified_purchase, review_headline, review_body, review_date, year,
product_category`
`from (`
  `select marketplace, customer_id, review_id, product_id, product_parent, product_title, star_rating,
helpful_votes, total_votes, vine, verified_purchase, review_headline, review_body, review_date, year,
product_category, ROW_NUMBER() over (partition by customer_id, product_id, product_category
order by review_date desc) as rank1`
`from amazon_review.amazon_reviews_parquet`
`) as temp`
`where temp.rank1 > 1;`
`select count(*) from amazon_reviews_exclude; --7159367`

```

```

hive> CREATE TABLE amazon_reviews_exclude AS
> select marketplace, customer_id, review_id, product_id, product_parent, product_title, star_rating, helpful_votes, total_votes, vine, verified_purchase,
review_headline, review_body, review_date, year, product_category
> from (
>   select marketplace, customer_id, review_id, product_id, product_parent, product_title, star_rating, helpful_votes, total_votes, vine, verified_purchas
e, review_headline, review_body, review_date, year, product_category, ROW_NUMBER() over (partition by customer_id, product_id, product_category order by revie
w_date desc) as rank1
>   from amazon_review.amazon_reviews_parquet
> ) as temp
> where temp.rank1 > 1;
Query ID = hadoop_20200413015201_069df257-46dc-4156-b4aa-163fef3e1e7c
Total jobs = 1
Launching Job 1 out of 1
tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1586737734097_0011)

-----
VERTICES    MODE      STATUS TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED 13      13          0         0        0        0
Reducer 2 ..... container  SUCCEEDED 39      39          0         0        0        0
-----
VERTICES: 02/02 [=====] 100% ELAPSED TIME: 461.51 s
-----
Moving data to directory hdfs:///ip-172-31-79-160.ec2.internal:8020/user/hive/warehouse/amazon_review.db/amazon_reviews_exclude
OK
Time taken: 475.633 seconds

```

6. Create a category view which will apply filters for our reviews

We created a view that will only consider those records with year  $\geq$  2005 and should be in one of the categories we need.

Query:

```
`CREATE OR REPLACE VIEW amazon_reviews_category AS`
`SELECT * FROM amazon_review.amazon_reviews_parquet`
`WHERE ("product_category" IN ('Wireless', 'Automotive ', 'Music', 'Digital_Music_Purchase', 'Sports',
'Toys', 'Digital_Video_Games', 'Video_Games'))`
`AND "year" >= 2005);`
`select count(*) from amazon_reviews_category;` -- 26898419
```

```
hive> CREATE OR REPLACE view amazon_reviews_category AS
> (SELECT * FROM amazon_review.amazon_reviews_parquet
> WHERE product_category IN ('Wireless', 'Automotive ', 'Music', 'Digital_Music_Purchase', 'Sports', 'Toys', 'Digital_Video_Games', 'Video_Games')
> AND year >= 2005);
OK
Time taken: 0.578 seconds
hive> select count(*) from amazon_review.amazon_reviews_category;
Query ID = hadoop_20200413021605_bbca2328-36dc-468b-9445-30b56a43b649
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1586737734097_0012)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED   13         13         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    1          1         0         0         0         0
-----
VERTICES: 02/02 [=====] 100% ELAPSED TIME: 31.98 s
-----
OK
26898419
Time taken: 32.739 seconds, Fetched: 1 row(s)
```

7. Create an include table that will be a join between the exclude table and the category view and will be our final table to query on.

Query:

```
`CREATE OR REPLACE VIEW amazon_reviews_include AS`
`select a.*`
`from amazon_review.amazon_reviews_parquet a`
`join amazon_review.amazon_reviews_category c`
`on a.review_id = c.review_id and a.product_category = c.product_category`
`where not exists`
`(select review_id, product_category from amazon_review.amazon_reviews_exclude b`
`where a.product_category = b.product_category`
`and a.review_id = b.review_id);`
```

```

hive> CREATE OR REPLACE VIEW amazon_reviews_include AS
> select a.*
> from amazon_review.amazon_reviews_parquet a
> join amazon_review.amazon_reviews_category c
> on a.review_id = c.review_id and a.product_category = c.product_category
> where not exists
> (select review_id, product_category from amazon_review.amazon_reviews_exclude b
>  where a.product_category = b.product_category
>  and a.review_id = b.review_id);
OK
Time taken: 0.456 seconds
hive> select count(*) from amazon_review.amazon_reviews_include;
Query ID = hadoop_20200413022924_887e056f-2115-40db-87d4-3224ae742b63
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1586737734097_0013)

```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1		container	INITED	13	0	0	13	0	0
Map 5		container	INITED	13	0	0	13	0	0
Map 6		container	INITED	14	0	0	14	0	0
Reducer 2		container	INITED	26	0	0	26	0	0
Reducer 3		container	INITED	22	0	0	22	0	0
Reducer 4		container	INITED	1	0	0	1	0	0
Reducer 7		container	INITED	1	0	0	1	0	0

```

VERTICES: 00/07  [>>-----] 0%  ELAPSED TIME: 5.42 s

```

```
`select count(*) from amazon_review.amazon_reviews_include;` -- 25335861
```

Using AWS EMR Hive and AWS Athena, answer the following questions:

### 1. Explore the dataset and provide basic exploratory analysis:

#### 1. Number of reviews

To calculate the number of reviews in our final table – amazon\_reviews\_include

Query:

```
`select count(review_id) from amazon_reviews_include;`
```

Output:

25335861

```

hive> select count(review_id) from amazon_reviews_include;
Query ID = hadoop_20200413023841_a4593dc0-00a5-4e36-875a-08f4000fd5dc
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1586737734097_0013)

```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	.....	container	SUCCEEDED	13	13	0	0	0	0
Map 5	.....	container	SUCCEEDED	13	13	0	0	0	0
Map 6	.....	container	SUCCEEDED	14	14	0	0	0	0
Reducer 2	.....	container	SUCCEEDED	26	26	0	0	0	0
Reducer 3	.....	container	SUCCEEDED	22	22	0	0	0	0
Reducer 4	.....	container	SUCCEEDED	1	1	0	0	0	0
Reducer 7	.....	container	SUCCEEDED	1	1	0	0	0	0

```

VERTICES: 07/07  [=====] 100%  ELAPSED TIME: 256.58 s
OK
25335861
Time taken: 257.804 seconds, Fetched: 1 row(s)
hive> █

```

## 2. Number of users

To calculate the number of users a.k.a number of customer ids in our final table.

Query:

```
`select count(customer_id) from amazon_reviews_include;`
```

Output:

25335861

```
hive> select count(customer_id) from amazon_reviews_include;
Query ID = hadoop_20200413025520_e1cd38c0-7034-4173-8cd9-9d0c1317de65
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1586737734097_0014)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED  13      13          0         0         0         0
Map 5 ..... container  SUCCEEDED  13      13          0         0         0         0
Map 6 ..... container  SUCCEEDED  14      14          0         0         0         0
Reducer 2 ..... container  SUCCEEDED  26      26          0         0         0         0
Reducer 3 ..... container  SUCCEEDED  22      22          0         0         0         0
Reducer 4 ..... container  SUCCEEDED   1         1          0         0         0         0
Reducer 7 ..... container  SUCCEEDED   1         1          0         0         0         0
-----
VERTICES: 07/07 [=====] 100% ELAPSED TIME: 280.76 s
-----
OK
25335861
Time taken: 290.833 seconds, Fetched: 1 row(s)
```

## 3. Average review stars

To calculate the average review stars for all years from 2005 for all product categories.

Query:

```
`select avg(star_rating) as avg_stars from amazon_reviews_include;`
```

Output:

4.153803614568299

```
hive> select avg(star_rating) as avg_stars from amazon_reviews_include;
Query ID = hadoop_20200413031617_47f31a25-06cf-4877-954c-0b27bfd9c880
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1586737734097_0015)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED  13      13          0         0         0         0
Map 5 ..... container  SUCCEEDED  13      13          0         0         0         0
Map 6 ..... container  SUCCEEDED  14      14          0         0         0         0
Reducer 2 ..... container  SUCCEEDED  26      26          0         0         0         0
Reducer 3 ..... container  SUCCEEDED  22      22          0         0         0         0
Reducer 4 ..... container  SUCCEEDED   1         1          0         0         0         0
Reducer 7 ..... container  SUCCEEDED   1         1          0         0         0         0
-----
VERTICES: 07/07 [=====] 100% ELAPSED TIME: 279.28 s
-----
OK
4.153803614568299
Time taken: 287.888 seconds, Fetched: 1 row(s)
```

## 4. Average length of the review

To calculate the average length of reviews by customers.

Query:

```
`select avg(length(review_body)) as avg_review_len from amazon_reviews_include;`
```

Output:

313.9593731020216

```
hive> select avg(length(review body)) as avg_review_len from amazon_reviews_include;
Query ID = hadoop_20200413032910_0acb63a2-ce35-42d4-a982-fd2b841eff42
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1586737734097_0016)
```

VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1 .....	container	SUCCEEDED	13	13	0	0	0	0
Map 5 .....	container	SUCCEEDED	13	13	0	0	0	0
Map 6 .....	container	SUCCEEDED	14	14	0	0	0	0
Reducer 2 .....	container	SUCCEEDED	26	26	0	0	0	0
Reducer 3 .....	container	SUCCEEDED	22	22	0	0	0	0
Reducer 4 .....	container	SUCCEEDED	1	1	0	0	0	0
Reducer 7 .....	container	SUCCEEDED	1	1	0	0	0	0

```
VERTICES: 07/07 [=====] 100% ELAPSED TIME: 590.46 s
OK
313.9593731020216
Time taken: 598.777 seconds, Fetched: 1 row(s)
```

## 5. Number of verified versus unverified reviews

To calculate the number of verified and unverified purchases of products on amazon by the customers.

Query:

```
` select 'verified' as review_type, count() from amazon_reviews_include where verified_purchase = 'Y'
union
select 'unverified' as review_type, count() from amazon_reviews_include where verified_purchase = 'N'`
```

Output:

Verified Purchase – 20624292

Unverified Purchase - 4711569



```
hive> select 'verified' as review_type, count(*) from amazon_review.amazon_reviews_include where verified_purchase = 'Y'
> union
> select 'unverified' as review_type, count(*) from amazon_review.amazon_reviews_include where verified_purchase = 'N';
No Stats for amazon_review@amazon_reviews_parquet, Columns: review_id, verified_purchase
No Stats for amazon_review@amazon_reviews_parquet, Columns: review_id, year
No Stats for amazon_review@amazon_reviews_exclude, Columns: review_id, product_category
No Stats for amazon_review@amazon_reviews_parquet, Columns: review_id, verified_purchase
No Stats for amazon_review@amazon_reviews_parquet, Columns: review_id, year
No Stats for amazon_review@amazon_reviews_exclude, Columns: review_id, product_category
Query ID = hadoop_20200413035137_f64e6ccd-c3fa-40b3-9d96-245fb729fe4c
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1586737734097_0017)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED  13      13          0        0        0        0
Map 10 ..... container  SUCCEEDED  13      13          0        0        0        0
Map 14 ..... container  SUCCEEDED  13      13          0        0        0        0
Map 15 ..... container  SUCCEEDED  14      14          0        0        0        0
Map 7 ..... container  SUCCEEDED  13      13          0        0        0        0
Map 8 ..... container  SUCCEEDED  14      14          0        0        0        0
Reducer 11 ..... container  SUCCEEDED  17      17          0        0        0        0
Reducer 12 ..... container  SUCCEEDED  8        8          0        0        0        0
Reducer 13 ..... container  SUCCEEDED  1        1          0        0        0        0
Reducer 16 ..... container  SUCCEEDED  1        1          0        0        0        0
Reducer 2 ..... container  SUCCEEDED  17      17          0        0        0        0
Reducer 3 ..... container  SUCCEEDED  8        8          0        0        0        0
Reducer 4 ..... container  SUCCEEDED  1        1          0        0        0        0
Reducer 6 ..... container  SUCCEEDED  1        1          0        0        0        0
Reducer 9 ..... container  SUCCEEDED  1        1          0        0        0        0
-----
VERTICES: 15/15 [=====] 100% ELAPSED TIME: 342.79 s
-----
OK
unverified      4711569
verified        20624292
Time taken: 355.036 seconds, Fetched: 2 row(s)
```

## 6. Top 10 Customer IDs with most helpful votes

To display the top 10 customers with the most number of helpful votes.

Query:

```
`select customer_id, helpful_votes from amazon_reviews_include`
```

```
`order by helpful_votes desc limit 10;`
```

Output:

Customer\_id      helpful\_votes

```
51394083      12188
9286343 10898
48557141      10498
33209578      9127
34072304      8650
9286343 8462
31076930      7624
9286343 7379
48475025      7166
15886460      6246
```

```
hive> select customer_id, helpful_votes from amazon_reviews_include
> order by helpful_votes desc limit 10;
Query ID = hadoop_20200413040339_750ea8c4-69b3-4ec6-a057-7287a952010b
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1586737734097_0018)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED  13      13          0        0        0        0
Map 5 ..... container  SUCCEEDED  13      13          0        0        0        0
Map 6 ..... container  SUCCEEDED  14      14          0        0        0        0
Reducer 2 ..... container  SUCCEEDED  26      26          0        0        0        0
Reducer 3 ..... container  SUCCEEDED  22      22          0        0        0        0
Reducer 4 ..... container  SUCCEEDED  1        1          0        0        0        0
Reducer 7 ..... container  SUCCEEDED  1        1          0        0        0        0
-----
VERTICES: 07/07 [=====] 100% ELAPSED TIME: 303.35 s
-----
OK
51394083      12188
9286343 10898
48557141      10498
33209578      9127
34072304      8650
9286343 8462
31076930      7624
9286343 7379
48475025      7166
15886460      6246
Time taken: 312.81 seconds, Fetched: 10 row(s)
hive> █
```

7. Most number of product categories reviewed in a given year

To calculate the total number of product categories reviewed in a given year.

Query:

```
`select year, count(product_category) as NoOfCategoriesReviewed`
`from amazon_reviews_include`
`group by year`
`order by NoOfCategoriesReviewed desc`
```

Output:

Year      NoOfCategoriesReviewed

```
2014      7341673
2015      7240577
2013      4870797
2012      2025258
2011      1122607
2010      710950
2009      563538
2008      445906
2007      394650
2005      334615
2006      285290
```

```
hive> select year, count(product_category) as NoOfCategoriesReviewed
> from amazon_reviews_include
> group by year
> order by NoOfCategoriesReviewed desc
> ;
Query ID = hadoop_20200413041308_001e4076-105a-43eb-80b7-37d55ac16858
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1586737734097_0018)

-----
VERTICES      MODE      STATUS      TOTAL      COMPLETED      RUNNING      PENDING      FAILED      KILLED
-----
Map 1 ..... container  SUCCEEDED      13           13           0           0           0           0
Map 6 ..... container  SUCCEEDED      13           13           0           0           0           0
Map 7 ..... container  SUCCEEDED      14           14           0           0           0           0
Reducer 2 ..... container  SUCCEEDED      26           26           0           0           0           0
Reducer 3 ..... container  SUCCEEDED      22           22           0           0           0           0
Reducer 4 ..... container  SUCCEEDED      12           12           0           0           0           0
Reducer 5 ..... container  SUCCEEDED       1            1           0           0           0           0
Reducer 8 ..... container  SUCCEEDED       1            1           0           0           0           0
-----
VERTICES: 08/08 [=====>>>] 100% ELAPSED TIME: 275.35 s
-----
OK
2014      7341673
2015      7240577
2013      4870797
2012      2025258
2011      1122607
2010      710950
2009      563538
2008      445906
2007      394650
2005      334615
2006      285290
Time taken: 276.457 seconds, Fetched: 11 row(s)
hive> █
```

8. Provide trending (over time) analysis of each of the metrics above

Queries:

-- marketplace

To calculate the number of reviews in a marketplace in a given year.

Query:

```
select year, marketplace, count(*) as CountMarketplace
from amazon_reviews_include
group by year, marketplace
order by CountMarketplace desc;
```

Output:

2014	US	7163174
2015	US	7122066
2013	US	4712678
2012	US	1956272
2011	US	1074584
2010	US	673787
2009	US	532768
2008	US	425332
2007	US	372071
2005	US	315417
2006	US	267167
2014	UK	105699
2015	UK	85514
2013	UK	73417
2013	DE	47926
2014	DE	38520
2012	UK	28449
2012	DE	21772
2013	FR	20262
2011	UK	20080
2014	FR	20078
2015	DE	16995
2013	JP	16514
2010	UK	14882
2014	JP	14202
2011	DE	14040
2009	UK	12562
2010	DE	11381
2012	FR	10068
2009	DE	9478
2012	JP	8697
2015	FR	8292
2007	UK	8011
2007	DE	7880
2008	UK	7843
2015	JP	7710
2011	JP	7135
2008	DE	6884
2011	FR	6768
2006	UK	6660
2005	UK	6549
2005	DE	6520



### -- customer\_id

To count the number of users reviewing products in a given year

Query:

```
select year, count(customer_id) as CountCustomerID
from amazon_reviews_include
group by year
order by CountCustomerID desc;
```

Output:

2014	7341673
2015	7240577
2013	4870797
2012	2025258
2011	1122607
2010	710950
2009	563538
2008	445906
2007	394650
2005	334615
2006	285290

### -- review\_id

To count the number of reviews reviewing products in a given year

Query:

```
select year, count(review_id) as CountReviewID
from amazon_reviews_include
group by year
```

order by CountReviewID desc;

Output:

2014	7341673
2015	7240577
2013	4870797
2012	2025258
2011	1122607
2010	710950
2009	563538
2008	445906
2007	394650
2005	334615
2006	285290

#### -- product\_id

To count the number of products in a given year.

Query:

```
select year, count(product_id) as CountProductID
from amazon_reviews_include
group by year
order by CountProductID desc;
```

Output:

2014	7341673
2015	7240577
2013	4870797
2012	2025258
2011	1122607
2010	710950
2009	563538
2008	445906
2007	394650
2005	334615
2006	285290

#### -- star\_rating

To calculate the average star rating for all the products in a given year ordered in decreasing order.

Query:

```
select year, avg(star_rating) as AvgStarRating
from amazon_reviews_include
group by year
order by AvgStarRating desc;
```

Output:

2015	4.205245659289308
2007	4.200709489421006
2014	4.174557624672197
2006	4.150162992043184
2005	4.144138786366421
2008	4.139654994550421
2013	4.134080315808686
2009	4.117454368649496
2012	4.071837760917375
2010	4.037252971376327
2011	4.0047282797987185

**-- helpful\_votes**

To calculate the average of helpful votes in a given year ordered by the average votes in descending order.

Query:

```
select year, avg(helpful_votes) as AvgHelpfulVotes
```


```
from amazon_reviews_include
```

```
group by year
```

```
order by AvgHelpfulVotes desc;
```

Output:

2005	5.532498543101774
2006	5.460384170493183
2007	4.217027746104143
2008	3.7557713957650267
2010	3.560725789436669
2009	3.5358414161955363
2011	2.829691958093972
2012	1.8801802042011437
2013	1.0723074683670866
2014	0.7410640599220368
2015	0.5606886025796011



**-- total\_votes**

To calculate average of total votes by customers in a given year and ordered by the average of total votes in descending votes.

```
select year, avg(total_votes) as AvgTotalVotes
```


```
from amazon_reviews_include
```

```
group by year
```

```
order by AvgTotalVotes desc;
```

Output:

2005	7.754326614168522
2006	7.562052648182551
2007	5.689045990117826
2008	5.02089902356102
2009	4.719846753901245
2010	4.677950629439483
2011	3.758173608395458
2012	2.524731663817647
2013	1.4987192034486347
2014	1.0520565271703057
2015	0.786011252970585

**-- verified\_purchase**

To calculate the number of verified and unverified purchases by all the customers starting from year 2005.

Query:

```
select 'verified' as review_type, count() from amazon_reviews_include where verified_purchase = 'Y'
```

```
union
```

```
select 'unverified' as review_type, count() from amazon_reviews_include where verified_purchase = 'N'
```

Output:

```

hive> select 'verified' as review_type, count(*) from amazon_review.amazon_reviews_include where verified_purchase = 'Y'
> union
> select 'unverified' as review_type, count(*) from amazon_review.amazon_reviews_include where verified_purchase = 'N';
No Stats for amazon_review@amazon_reviews_parquet, Columns: review_id, verified_purchase
No Stats for amazon_review@amazon_reviews_parquet, Columns: review_id, year
No Stats for amazon_review@amazon_reviews_exclude, Columns: review_id, product_category
No Stats for amazon_review@amazon_reviews_parquet, Columns: review_id, verified_purchase
No Stats for amazon_review@amazon_reviews_parquet, Columns: review_id, year
No Stats for amazon_review@amazon_reviews_exclude, Columns: review_id, product_category
Query ID = hadoop_20200413035137_f64e6ccd-c3fa-40b3-9d96-245fb729fe4c
Total jobs = 1
Launching Job 1 out of 1
Tez session was closed. Reopening...
Session re-established.
Status: Running (Executing on YARN cluster with App id application_1586737734097_0017)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED  13      13          0         0         0         0
Map 10 ..... container  SUCCEEDED  13      13          0         0         0         0
Map 14 ..... container  SUCCEEDED  13      13          0         0         0         0
Map 15 ..... container  SUCCEEDED  14      14          0         0         0         0
Map 7 ..... container  SUCCEEDED  13      13          0         0         0         0
Map 8 ..... container  SUCCEEDED  14      14          0         0         0         0
Reducer 11 ..... container  SUCCEEDED  17      17          0         0         0         0
Reducer 12 ..... container  SUCCEEDED  8        8          0         0         0         0
Reducer 13 ..... container  SUCCEEDED  1        1          0         0         0         0
Reducer 16 ..... container  SUCCEEDED  1        1          0         0         0         0
Reducer 2 ..... container  SUCCEEDED  17      17          0         0         0         0
Reducer 3 ..... container  SUCCEEDED  8        8          0         0         0         0
Reducer 4 ..... container  SUCCEEDED  1        1          0         0         0         0
Reducer 6 ..... container  SUCCEEDED  1        1          0         0         0         0
Reducer 9 ..... container  SUCCEEDED  1        1          0         0         0         0
-----
VERTICES: 15/15 [=====] 100% ELAPSED TIME: 342.79 s
-----
OK
unverified      4711569
verified        20624292
Time taken: 355.036 seconds, Fetched: 2 row(s)

```

### -- product\_category

To calculate the number of product categories reviewed in a given year.

Query:

```
select year, product_category, count(*) as record_count from amazon_reviews_include
```

```
group by year, product_category
```

```
order by record_count desc;
```

Output:


```

2014    Digital_Video_Games    43751
2007    Video_Games             43485
2013    Digital_Video_Games    43263
2008    Sports                    40925
2010    Digital_Music_Purchase    40809
2009    Digital_Music_Purchase    36182
2005    Toys                      36105
2015    Digital_Video_Games    30022
2007    Sports                    29541
2005    Video_Games             27102
2006    Toys                      26849
2006    Video_Games             24702
2008    Digital_Music_Purchase    22040
2006    Wireless                 19855
2012    Digital_Video_Games    16624
2005    Wireless                 11834
2006    Sports                    9529
2011    Digital_Video_Games    7644
2005    Sports                    4514
2010    Digital_Video_Games    2551
2007    Digital_Music_Purchase    2235
2009    Digital_Video_Games    1561
2006    Digital_Music_Purchase    21
2005    Digital_Music_Purchase    8
2008    Digital_Video_Games    5
2006    Digital_Video_Games    1

```

Only displaying the last 20 rows





2. Provide detailed analysis of Music/Digital\_Music\_Purchase and Digital\_Video\_Games/Video\_Games over time.

Query:

-- Music/Digital\_Music\_Purchase over time.

We created two separate tables for product categories 'Music' and 'Digital Music Purchase' as it becomes easier to compare the two categories and answer questions based on it.

Query:

```
`create table amazon_review.sample_music as `
```

```
`(select * from amazon_reviews_include where product_category = 'Music') with data;`
```

```
`create table amazon_review.dmp as`
```

```
`(select * from amazon_reviews_include where product_category = 'Digital_Music_Purchase') with data;`
```

Output:

```

-----
      VERTICES      MODE      STATUS      TOTAL      COMPLETED      RUNNING      PENDING      FAILED      KILLED
-----
Map 1 ..... container  SUCCEEDED      20          20           0           0           0           0
Map 4 ..... container  SUCCEEDED      20          20           0           0           0           0
Map 5 ..... container  SUCCEEDED      14          14           0           0           0           0
Reducer 2 ..... container  SUCCEEDED      16          16           0           0           0           0
Reducer 3 ..... container  SUCCEEDED       5           5           0           0           0           0
Reducer 6 ..... container  SUCCEEDED       1           1           0           0           0           0
-----
VERTICES: 06/06 [=====>>>] 100% ELAPSED TIME: 256.81 s
-----
Moving data to directory hdfs://ip-172-31-79-160.ec2.internal:8020/user/hive/warehouse/amazon_review.db/sample_music
OK
Time taken: 266.291 seconds
hive> create table amazon_review.dmp as
> (select * from amazon_reviews_include where product_category = 'Digital_Music_Purchase');
Query ID = hadoop_20200413055455_2e3e619a-81d1-429d-b08f-6e17662ff015
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1586737734097_0021)
-----
      VERTICES      MODE      STATUS      TOTAL      COMPLETED      RUNNING      PENDING      FAILED      KILLED
-----
Map 1 ..... container  SUCCEEDED      10          10           0           0           0           0
Map 4 ..... container  SUCCEEDED      10          10           0           0           0           0
Map 5 ..... container  SUCCEEDED      14          14           0           0           0           0
Reducer 2 ..... container  SUCCEEDED       3           3           0           0           0           0
Reducer 3 ..... container  SUCCEEDED       1           1           0           0           0           0
Reducer 6 ..... container  SUCCEEDED       1           1           0           0           0           0
-----
VERTICES: 06/06 [=====>>>] 100% ELAPSED TIME: 113.53 s
-----
Moving data to directory hdfs://ip-172-31-79-160.ec2.internal:8020/user/hive/warehouse/amazon_review.db/dmp
OK

```

1. Do you see correlation (maybe negative) between the categories over time?

We attempt to correlate the star ratings of both product categories and use `corr()` function to determine the correlation coefficient.

Query:

```

select corr(b.star_rating, c.star_rating)

from(

select year, avg(star_rating) as star_rating

from amazon_review.dmp

group by year

) as b

join

(select year, avg(star_rating) as star_rating

from amazon_review.sample_music

group by year

) as c

on b.year = c.year;

```

Output:

0.5806308546262009

```
hive> select corr(b.star_rating, c.star_rating)
> from
> (
> select year, avg(star_rating) as star_rating
> from amazon_review.dmp
> group by year
> ) as b
> join
> (
> select year, avg(star_rating) as star_rating
> from amazon_review.sample_music
> group by year
> ) as c
> on b.year = c.year;
Query ID = hadoop_20200413060056_f44b716b-17b6-4661-8ec3-deb9ef2c390c
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1586737734097_0021)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED   11         11         0         0         0         0
Map 5 ..... container  SUCCEEDED   17         17         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    3          3         0         0         0         0
Reducer 3 ..... container  SUCCEEDED    6          6         0         0         0         0
Reducer 4 ..... container  SUCCEEDED    1          1         0         0         0         0
Reducer 6 ..... container  SUCCEEDED   10         10         0         0         0         0
-----
VERTICES: 06/06 [=====] 100% ELAPSED TIME: 39.87 s
-----
OK
0.5806308546262009
Time taken: 40.532 seconds, Fetched: 1 row(s)
hive> █
```

Based on the results, we can say that they are positively correlated.

2. Are there same users reviewing in both categories?

We create a new table called 'final' which contains all the records with product category as Music or Digital Music Purchase and which contains the same customer ids in both tables. We attempt to get the number of users who reviewed products for both the categories.

Query:

```
`create table amazon_review.final as`
`(select a.*`
`from sample_music a`
`join dmp b`
`on a.customer_id = b.customer_id);`
```

Select count(distinct(customer\_id)) from final; -- 140825

Output:

There are 140825 distinct users who have reviewed products from both categories.

```
hive> select count(distinct(customer_id)) from final;
Query ID = hadoop_20200413063804_dd87051e-a6e3-49bd-85e9-1a835468f24a
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1586737734097_0022)
```

VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1 .....	container	SUCCEEDED	30	30	0	0	0	0
Reducer 2 .....	container	SUCCEEDED	107	107	0	0	0	0
Reducer 3 .....	container	SUCCEEDED	1	1	0	0	0	0

```
VERTICES: 03/03 [=====>>] 100% ELAPSED TIME: 270.40 s
OK
140825
Time taken: 271.266 seconds, Fetched: 1 row(s)
hive>
```

```
hive> select count(distinct(customer_id)) from final;
Query ID = hadoop_20200413063804_dd87051e-a6e3-49bd-85e9-1a835468f24a
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1586737734097_0022)
```

VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1 .....	container	SUCCEEDED	30	30	0	0	0	0
Reducer 2 .....	container	SUCCEEDED	107	107	0	0	0	0
Reducer 3 .....	container	SUCCEEDED	1	1	0	0	0	0

```
VERTICES: 03/03 [=====>>] 100% ELAPSED TIME: 270.40 s
OK
140825
Time taken: 271.266 seconds, Fetched: 1 row(s)
hive>
```

3. Can you identify similar items in both categories? Do they get same rating?

To identify similar items for both categories, we compare their product ids and check if we have anything in common. We then display their star ratings and see that the star ratings for these products are different.

Query:

```
`select product_id from dmp`
```

```
`Intersect`
```

```
`select product_id from sample_music`
```

```
`select star_rating, product_id from sample_music where product_id = 'B0019M1ZJS'`
```

```
`union`
```

```
`select star_rating, product_id from dmp where product_id = 'B0019M1ZJS'`
```

Output:

```

Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1586737734097_0022)

-----
      VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    6         6         0         0         0         0
Map 5 ..... container  SUCCEEDED    7         7         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    3         3         0         0         0         0
Reducer 4 ..... container  SUCCEEDED    6         6         0         0         0         0
Reducer 6 ..... container  SUCCEEDED   10        10         0         0         0         0
-----
VERTICES: 05/05  [=====>>>] 100%  ELAPSED TIME: 79.60 s
-----
OK
B0019M1ZJS
Time taken: 80.726 seconds, Fetched: 1 row(s)
hive> select star_rating, product_id from sample_music where product_id = 'B0019M1ZJS'
> union
> select star_rating, product_id from dmp where product_id = 'B0019M1ZJS'
> ;
Query ID = hadoop_20200413064829_5d8f7975-c474-4843-a982-cdd26b0bd2b3
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1586737734097_0022)

-----
      VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    7         7         0         0         0         0
Map 4 ..... container  SUCCEEDED    6         6         0         0         0         0
Reducer 3 ..... container  SUCCEEDED    6         6         0         0         0         0
-----
VERTICES: 03/03  [=====>>>] 100%  ELAPSED TIME: 38.42 s
-----
OK
5      B0019M1ZJS
3      B0019M1ZJS
Time taken: 39.245 seconds, Fetched: 2 row(s)
hive> █

```

#### 4. Comparing average star ratings for both the categories grouped by years

To compare the average star ratings for both the categories in a given year starting from 2005.

Query:

```

`select year, product_category, avg(star_rating) as AvgStarRating`
`from sample_music`
`group by year, product_category`
`union`
`select year, product_category, avg(star_rating) as AvgStarRtnng`
`from dmp`
`group by year, product_category`


```

Output:

```

2006 Digital_Music_Purchase 3.857142857142857
2007 Digital_Music_Purchase 4.405369127516779
2008 Digital_Music_Purchase 4.4610254083484575
2009 Digital_Music_Purchase 4.461279089049804
2015 Music 4.619308172689674
2007 Music 4.38456067207644
2010 Digital_Music_Purchase 4.49207282707246
2011 Digital_Music_Purchase 4.479447986723732
2011 Music 4.419066943157351
2014 Digital_Music_Purchase 4.672041015885251
2005 Digital_Music_Purchase 4.625
2010 Music 4.405982861158602
2012 Digital_Music_Purchase 4.612081540274397
2013 Music 4.538666302915294
2015 Digital_Music_Purchase 4.703664618181714
2012 Music 4.474767487751498
2013 Digital_Music_Purchase 4.6466514195021045
2014 Music 4.579786644047891
2006 Music 4.3171049218677355
2009 Music 4.390710809701566
2005 Music 4.271732823110582
2008 Music 4.3734695457007975
Time taken: 41.238 seconds, Fetched: 22 row(s)
hive>

```



5. Comparing number of 5-star ratings with 4-star ratings for reviews for both categories grouped by year

To compare the no of 5-star and 4-star ratings for reviews from both categories in a given year. We then try to plot the results to find a trend in the way the star ratings change.

Query:

```

`select year, star_rating, product_category, count(*) as count_star_rating`

`from sample_music`

`where star_rating = 5`

`group by year, star_rating, product_category`

`union`

`select year, star_rating, product_category, count(*) as count_star_rating`

```

```

`from sample_music`

`where star_rating = 4`

`group by year, star_rating, product_category`

`union`

`select year, star_rating, product_category, count(*) as count_star_rating`

`from dmp`

`where star_rating = 5`

`group by year, star_rating, product_category`

`union`

`select year, star_rating, product_category, count(*) as count_star_rating`

`from dmp`

`where star_rating = 4`

`group by year, star_rating, product_category`

```


Output:

```

2008 4 Music 37781
2010 4 Music 34246
2010 5 Digital_Music_Purchase 30282
2011 4 Digital_Music_Purchase 7362
2011 4 Music 35415
2012 4 Music 40215
2014 5 Music 466682
2015 4 Music 60756
2015 5 Music 390166
2005 4 Digital_Music_Purchase 1
2006 5 Digital_Music_Purchase 11
2007 4 Digital_Music_Purchase 341
2007 4 Music 44083
2008 5 Music 125744
2009 4 Digital_Music_Purchase 5197
2009 4 Music 38549
2009 5 Music 135065
2010 4 Digital_Music_Purchase 5486
2012 5 Digital_Music_Purchase 150911
2012 5 Music 177683
2013 4 Digital_Music_Purchase 58772
2013 4 Music 75527
2013 5 Digital_Music_Purchase 354514
2013 5 Music 365376
2014 5 Digital_Music_Purchase 395154
2005 4 Music 50017
2006 4 Music 40115
2007 5 Music 145545
2008 5 Digital_Music_Purchase 15979
2009 5 Digital_Music_Purchase 26320
2010 5 Music 127987
2011 5 Digital_Music_Purchase 42635
2011 5 Music 136812
2012 4 Digital_Music_Purchase 24866
2014 4 Digital_Music_Purchase 56924
2014 4 Music 81892
2015 4 Digital_Music_Purchase 37354
2015 5 Digital_Music_Purchase 288437
Time taken: 63.643 seconds, Fetched: 44 row(s)
hive> █

```

*Displaying only the last 20 lines*



#### -- Digital\_Video\_Games/Video\_Games over time.

We created two separate tables for product categories 'Video\_Games' and 'Digital\_Video\_Games' as it becomes easier to compare the two categories and answer questions based on it.

Query:

```
create table amazon_review.dvg as
```

```
(select * from amazon_reviews_include where product_category = 'Digital_Video_Games') with data;
```

```
create table amazon_review.sample_video as
```

```
(select * from amazon_reviews_include where product_category = 'Video_Games') with data;
```

Output:



```

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    2         2         0         0         0         0
Map 4 ..... container  SUCCEEDED    2         2         0         0         0         0
Map 5 ..... container  SUCCEEDED   14        14         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    1         1         0         0         0         0
Reducer 3 ..... container  SUCCEEDED    1         1         0         0         0         0
Reducer 6 ..... container  SUCCEEDED    1         1         0         0         0         0
-----
VERTICES: 06/06 [=====>>>] 100% ELAPSED TIME: 47.92 s
-----
Moving data to directory hdfs://ip-172-31-79-160.ec2.internal:8020/user/hive/warehouse/amazon_review.db/dvg
OK
Time taken: 49.867 seconds
hive> create table amazon_review.sample_video as
> (select * from amazon_reviews_include where product_category = 'Video_Games')
> ;
Query ID = hadoop_20200413070815_7ac2ec6c-203c-4759-b2d9-d3663b00309a
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1586737734097_0023)
-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED   11        11         0         0         0         0
Map 4 ..... container  SUCCEEDED   11        11         0         0         0         0
Map 5 ..... container  SUCCEEDED   14        14         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    4         4         0         0         0         0
Reducer 3 ..... container  SUCCEEDED    2         2         0         0         0         2
Reducer 6 ..... container  SUCCEEDED    1         1         0         0         0         0
-----
VERTICES: 06/06 [=====>>>] 100% ELAPSED TIME: 146.26 s
-----
Moving data to directory hdfs://ip-172-31-79-160.ec2.internal:8020/user/hive/warehouse/amazon_review.db/sample_video
OK
Time taken: 147.791 seconds
hive> █

```

1. Do you see correlation (maybe negative) between the categories over time?

We attempt to correlate the star ratings of both product categories and use `corr()` function to determine the correlation coefficient.

Query:

```

select corr(b.star_rating, c.star_rating)
from(
select year, avg(star_rating) as star_rating
from amazon_review.dvg
group by year
) as b
join
(select year, avg(star_rating) as star_rating
from amazon_review.sample_video
group by year
) as c
on b.year = c.year;

```

Output:

0.435964269159388

```
hive> select corr(b.star_rating, c.star_rating)
> from(
> select year, avg(star_rating) as star_rating
> from amazon_review.dvg
> group by year
> ) as b
> join
> (select year, avg(star_rating) as star_rating
> from amazon_review.sample_video
> group by year
> ) as c
> on b.year = c.year;
Query ID = hadoop_20200413071152_04b14140-1225-4c38-8b8d-a2f26a0e36cf
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1586737734097_0023)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    2         2         0         0         0         0
Map 5 ..... container  SUCCEEDED    8         8         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    1         1         0         0         0         0
Reducer 3 ..... container  SUCCEEDED    3         3         0         0         0         0
Reducer 4 ..... container  SUCCEEDED    1         1         0         0         0         0
Reducer 6 ..... container  SUCCEEDED    5         5         0         0         0         0
-----
VERTICES: 06/06 [=====] 100% ELAPSED TIME: 29.32 s
-----
OK
0.435964269159388
Time taken: 30.489 seconds, Fetched: 1 row(s)
hive>
```

After observing the results, we can say that the star ratings for both the categories are positively correlated.

2. Are there same users reviewing in both categories?

We create a new table called 'final2' which contains all the records with product category as Video\_Games or Digital\_Video\_Games and which contains the same customer ids in both tables. We attempt to get the number of users who reviewed products for both the categories.

Query:

```
create table amazon_review.final2 as
```

```
(select a.*
```

```
from dvg a
```

```
join sample_video b
```

```
on a.customer_id = b.customer_id);
```

```
select count(distinct(customer_id)) from final2 -- 29767
```

Output:

Number of users who reviewed the product for both categories are 29767

```
hive> create table amazon_review.final2 as
> (select a.*
> from dvg a
> join sample_video b
> on a.customer_id = b.customer_id)
> ;
Query ID = hadoop_20200413071506_7d970bc6-19f1-4d90-b9d9-57a122e16943
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1586737734097_0023)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    2         2         0         0         0         0
Map 3 ..... container  SUCCEEDED    8         8         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    5         5         0         0         0         0
-----
VERTICES: 03/03 [=====>>>] 100% ELAPSED TIME: 50.92 s
-----
Moving data to directory hdfs://ip-172-31-79-160.ec2.internal:8020/user/hive/warehouse/amazon_review.db/final2
OK
Time taken: 52.017 seconds
hive> select count(distinct(customer_id)) from final2;
Query ID = hadoop_20200413071622_77a26abd-865b-4e06-af40-73de05af9488
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1586737734097_0023)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    7         7         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    3         3         0         0         0         0
Reducer 3 ..... container  SUCCEEDED    1         1         0         0         0         0
-----
VERTICES: 03/03 [=====>>>] 100% ELAPSED TIME: 22.62 s
-----
OK
29767
Time taken: 23.171 seconds, Fetched: 1 row(s)
hive>
```

3. Can you identify similar items in both categories? Do they get same rating?

To identify similar items for both categories, we compare their product ids and check if we have anything in common. We then display their star ratings and see that the star ratings for these products are different.

Query:

```
select product_id from sample_video
```

```
intersect
```

```
select product_id from dvg
```

```
hive> select product_id from sample_video
> intersect
> select product_id from dvg
> ;
Query ID = hadoop_20200413071828_55ada534-f8b3-445e-b5b2-bd21c6bf6935
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1586737734097_0023)

-----
VERTICES      MODE      STATUS  TOTAL  COMPLETED  RUNNING  PENDING  FAILED  KILLED
-----
Map 1 ..... container  SUCCEEDED    8         8         0         0         0         0
Map 5 ..... container  SUCCEEDED    2         2         0         0         0         0
Reducer 2 ..... container  SUCCEEDED    5         5         0         0         0         0
Reducer 4 ..... container  SUCCEEDED    3         3         0         0         0         0
Reducer 6 ..... container  SUCCEEDED    1         1         0         0         0         0
-----
VERTICES: 05/05 [=====>>>] 100% ELAPSED TIME: 30.68 s
-----
OK
B004YNI19Y
B00B4WVTUS
B00NBME0Y
B0047TMEW
Time taken: 31.375 seconds, Fetched: 4 row(s)
hive>
```

```
select star_rating, product_category, product_id
from sample_video a
where product_id = 'B00B4WVTUS'
union
select star_rating, product_category, product_id
from dvg
where product_id = 'B00B4WVTUS'
union
select star_rating, product_category, product_id
from sample_video a
where product_id = 'B0047T7MEW'
union
select star_rating, product_category, product_id
from dvg
where product_id = 'B0047T7MEW'
union
select star_rating, product_category, product_id
from sample_video a
where product_id = 'B00NBBME0Y'
union
select star_rating, product_category, product_id
from dvg
where product_id = 'B00NBBME0Y'
union
select star_rating, product_category, product_id
from sample_video a
where product_id = 'B004YNII9Y'
union
select star_rating, product_category, product_id
from dvg
where product_id = 'B004YNII9Y'
```

Output:

```

1 Digital_Video_Games B004YNII9Y
2 Digital_Video_Games B0047T7MEW
3 Digital_Video_Games B004YNII9Y
3 Video_Games B00NBBME0Y
4 Digital_Video_Games B004YNII9Y
4 Video_Games B00NBBME0Y
5 Digital_Video_Games B0047T7MEW
5 Digital_Video_Games B004YNII9Y
5 Video_Games B0047T7MEW
5 Video_Games B00NBBME0Y
1 Digital_Video_Games B0047T7MEW
1 Digital_Video_Games B00B4WVTUS
1 Video_Games B00NBBME0Y
2 Digital_Video_Games B004YNII9Y
3 Digital_Video_Games B0047T7MEW
4 Digital_Video_Games B0047T7MEW
4 Video_Games B0047T7MEW
4 Video_Games B004YNII9Y
5 Digital_Video_Games B00B4WVTUS
5 Digital_Video_Games B00NBBME0Y
5 Video_Games B00B4WVTUS
Time taken: 48.284 seconds, Fetched: 21 row(s)
hive>

```



#### 4. Comparing average star ratings for both the categories grouped by years

To compare the average star ratings for both the categories in a given year starting from 2005.

Query:

```

select year, product_category, avg(star_rating) as AvgStarRating
from dvg
group by year, product_category
union
select year, product_category, avg(star_rating) as AvgStarRtnng
from sample_video
group by year, product_category

```


Output:

```

2007    Video_Games      3.955317925721513
2008    Video_Games      3.7906369966596314
2010    Digital_Video_Games  3.733829870638965
2011    Digital_Video_Games  3.6483516483516483
2005    Video_Games      3.765995129510737
2006    Digital_Video_Games  4.0
2006    Video_Games      3.751882438668934
2009    Digital_Video_Games  3.89237668161435
2009    Video_Games      3.9204883414587774
2010    Video_Games      3.8642923005993546
2011    Video_Games      3.8867357309397126
2012    Video_Games      3.949663479046681
2014    Digital_Video_Games  3.9444584123791455
2014    Video_Games      4.167348550380167
2015    Digital_Video_Games  4.027213376856971
2015    Video_Games      4.2352707879516265
2008    Digital_Video_Games  2.0
2012    Digital_Video_Games  3.683890760346487
2013    Digital_Video_Games  3.746943115364168
2013    Video_Games      4.097030029181964
Time taken: 26.343 seconds, Fetched: 20 row(s)
hive>

```

Sheet 8



5. Comparing number of 5-star ratings with 4-star ratings for reviews for both categories grouped by year

To compare the number of 5-star and 4-star ratings for reviews from both categories in a given year. We then try to plot the results to find a trend in the way the star ratings change.

Query:

```

select year, star_rating, product_category, count(*) as count_star_rating
from dvrg
where star_rating = 5
group by year, star_rating, product_category

```

union

```
select year, star_rating, product_category, count(*) as count_star_rating
```

```
from dvg
```

```
where star_rating = 4
```

```
group by year, star_rating, product_category
```

union

```
select year, star_rating, product_category, count(*) as count_star_rating
```

```
from sample_video
```

```
where star_rating = 5
```

```
group by year, star_rating, product_category
```

union

```
select year, star_rating, product_category, count(*) as count_star_rating
```


```
from sample_video
```

```
where star_rating = 4
```

```
group by year, star_rating, product_category
```

Output:

```
2005 5 Video_Games 12289
2007 4 Video_Games 10492
2007 5 Video_Games 21098
2008 4 Digital_Video_Games 1
2010 5 Digital_Video_Games 1102
2010 5 Video_Games 41698
2011 5 Digital_Video_Games 3389
2012 4 Digital_Video_Games 2962
2012 4 Video_Games 28316
2013 4 Digital_Video_Games 6209
2013 4 Video_Games 50371
2013 5 Digital_Video_Games 22630
2013 5 Video_Games 176534
2014 4 Video_Games 57169
2014 5 Digital_Video_Games 25789
2014 5 Video_Games 242720
2015 5 Video_Games 221051
2005 4 Video_Games 5973
2006 4 Digital_Video_Games 1
2006 5 Video_Games 5720
2006 5 Video_Games 10764
2008 4 Video_Games 13837
2008 5 Video_Games 27414
2009 4 Digital_Video_Games 344
2009 5 Video_Games 17171
2009 5 Digital_Video_Games 754
2009 5 Video_Games 35787
2010 4 Digital_Video_Games 592
2010 4 Video_Games 18976
2011 4 Digital_Video_Games 1510
2011 4 Video_Games 24650
2011 5 Video_Games 58038
2012 5 Digital_Video_Games 7896
2012 5 Video_Games 79297
2014 4 Digital_Video_Games 5601
2015 4 Digital_Video_Games 3185
2015 4 Video_Games 44683
2015 5 Digital_Video_Games 19111
Time taken: 35.101 seconds, Fetched: 38 row(s)
hive>
```



### 3. You should demonstrate your ability to use Hive advanced functions:

1. Window functions: moving average, rank, aggregation functions using relevant ordering and partitioning

Query:

#### -- moving average

To calculate a 10 day moving average over average star ratings from sample\_music table. We order the results by review\_date.

Query:

```
select review_date, star_rating, avg(total_votes) over (order by review_date asc rows 9 PRECEDING) as MA10
from sample_music
```

Output:

```
2015-08-31    5    0.3
2015-08-31    5    0.1
2015-08-31    5    0.2
2015-08-31    5    0.3
2015-08-31    3    0.3
2015-08-31    5    0.3
2015-08-31    4    0.3
2015-08-31    4    0.3
2015-08-31    5    0.3
2015-08-31    3    0.4
2015-08-31    5    0.4
2015-08-31    5    0.4
2015-08-31    1    0.4
2015-08-31    5    0.3
2015-08-31    5    0.3
2015-08-31    5    0.3
2015-08-31    5    0.4
2015-08-31    5    0.2
2015-08-31    5    0.2
2015-08-31    5    0.2
2015-08-31    5    0.1
2015-08-31    5    0.2
2015-08-31    5    0.2
Time taken: 43.749 seconds, Fetcher
hive>
```



**-- rank**

To use rank() function to rank the average star ratings grouped by years. We display the first ten ranks.

Query:

```
select a.* from
(select year, avg(star_rating) as AvgStarRtnng, RANK() over (order by avg(star_rating) desc) as rank1
from amazon_reviews_include
group by year ) as a
where rank1 <= 10;
```

Output:

2015	4.205245772675726	1
2007	4.200702414105042	2
2014	4.174557594447047	3
2006	4.150159486837954	4
2005	4.144136228609682	5
2008	4.139655620940831	6
2013	4.134083615611352	7
2009	4.117450819643041	8
2012	4.071836456975992	9
2010	4.037237656252857	10

**-- aggregate functions and partitioning**

Use aggregate functions and partitioning to choose such customers with their most recent review and exclude multiple reviews by that customer. Also, apply basic filters based on our requirement like starting our analysis from 2005 and over a few selected product categories.

Query:

```
select count(*) from
(select *, ROW_NUMBER() over (partition by customer_id, review_id order by review_date desc) as rank1 from
(select * from amazon_review.amazon_reviews_parquet
where year >= 2005
and product_category in
('Wireless','Automotive','Music','Digital_Music_Purchase','Sports','Toys','Digital_Video_Games','Video_Games')
) as a
) as b
where b.rank1 = 1
```

Output:

```

hive> select count(*) from
> (select *, ROW_NUMBER() over (partition by customer_id, review_id order by review_date desc) as rank1 from
> (select * from amazon_review.amazon_reviews_parquet
> where year >= 2005
> and product_category in ('Wireless','Automotive','Music','Digital_Music_Purchase','Sports','Toys','Digital_Video_Games','Video_Games')
> ) as a
> ) as b
> where b.rank1 = 1
> ;
Query ID = hadoop_20200413092404_8684d6f7-e6f7-4da0-a228-f518d416e4a7
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_158676680493_0011)

-----
VERTICES      MODE        STATUS      TOTAL   COMPLETED   RUNNING   PENDING   FAILED   KILLED
-----
Map 1 ..... container  SUCCEEDED    26         26           0         0         0         0
Reducer 2 ..... container  SUCCEEDED    13          13           0         0         0         0
Reducer 3 ..... container  SUCCEEDED     1           1           0         0         0         0
-----
VERTICES: 03/03 [=====] 100% ELAPSED TIME: 78.29 s
-----
OK
29635938
Time taken: 82.612 seconds, Fetched: 1 row(s)
hive>

```

## 2. Analytical Aggregate functions: percentile, min, max, average, standard deviation, correlation

### -- percentile

To calculate the percentile rank of star\_rating for all the categories grouped by year.

Query:

```

select year, star_rating, PERCENT_RANK() over (order by star_rating) as percent_rank
from amazon_reviews_include
group by year, star_rating

```

Output:

```

2007 3 0.4074074074074074
2008 3 0.4074074074074074
2006 3 0.4074074074074074
2005 3 0.4074074074074074
2009 3 0.4074074074074074
2011 3 0.4074074074074074
2013 3 0.4074074074074074
2014 3 0.4074074074074074
2012 3 0.4074074074074074
2006 4 0.6111111111111112
2011 4 0.6111111111111112
2012 4 0.6111111111111112
2013 4 0.6111111111111112
2010 4 0.6111111111111112
2014 4 0.6111111111111112
2005 4 0.6111111111111112
2008 4 0.6111111111111112
2015 4 0.6111111111111112
2007 4 0.6111111111111112
2009 4 0.6111111111111112
2006 5 0.8148148148148148
2007 5 0.8148148148148148
2005 5 0.8148148148148148
2015 5 0.8148148148148148
2011 5 0.8148148148148148
2013 5 0.8148148148148148
2009 5 0.8148148148148148
2010 5 0.8148148148148148
2012 5 0.8148148148148148
2014 5 0.8148148148148148
2008 5 0.8148148148148148
Time taken: 206.675 seconds, Fetched: 1 row(s)
hive>

```

-- min

To calculate minimum review\_date.

Query:

```
select min(review_date) as Min
```

```
from amazon_reviews_include
```

Output:

2005-01-01

```
hive> select min(review_date) as Min
> from amazon_reviews_include
> ;
Query ID = hadoop_20200413093104_94902787-ab0f-4dfa-b1cf-4c6e905b0fce
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1586766680493_0011)
```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	.....	container	SUCCEEDED	26	26	0	0	0	0
Map 5	.....	container	SUCCEEDED	26	26	0	0	0	0
Map 6	.....	container	SUCCEEDED	20	20	0	0	0	0
Reducer 2	.....	container	SUCCEEDED	26	26	0	0	0	0
Reducer 3	.....	container	SUCCEEDED	22	22	0	0	0	0
Reducer 4	.....	container	SUCCEEDED	1	1	0	0	0	0
Reducer 7	.....	container	SUCCEEDED	1	1	0	0	0	0

```
VERTICES: 07/07 [=====>>>] 100% ELAPSED TIME: 192.46 s
OK
2005-01-01
Time taken: 193.303 seconds, Fetched: 1 row(s)
hive>
```

-- max

To calculate maximum review\_date

Query:

```
select max(review_date) as Max
```

```
from amazon_reviews_include
```

Output:

```
hive> select max(review_date) as Max
> from amazon_reviews_include
> ;
Query ID = hadoop_20200413093716_df52906b-450d-44b3-ae3f-78af50eeaea1
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1586766680493_0011)
```

	VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	.....	container	SUCCEEDED	26	26	0	0	0	0
Map 5	.....	container	SUCCEEDED	26	26	0	0	0	0
Map 6	.....	container	SUCCEEDED	20	20	0	0	0	0
Reducer 2	.....	container	SUCCEEDED	26	26	0	0	0	0
Reducer 3	.....	container	SUCCEEDED	22	22	0	0	0	0
Reducer 4	.....	container	SUCCEEDED	1	1	0	0	0	0
Reducer 7	.....	container	SUCCEEDED	1	1	0	0	0	0

```
VERTICES: 07/07 [=====] 100% ELAPSED TIME: 197.47 s
OK
2015-08-31
Time taken: 198.22 seconds, Fetched: 1 row(s)
hive>
```

### -- average

To calculate average star rating grouped by year.

Query:

```
select year, avg(star_rating) as AVG_Rating
from amazon_reviews_include
group by year
order by AVG_Rating desc
```

Output:

```
2015 4.205245772675726
2007 4.200702414105042
2014 4.174557594447047
2006 4.150159486837954
2005 4.144136228609682
2008 4.139655620940831
2013 4.134083615611352
2009 4.117450819643041
2012 4.071836456975992
2010 4.037237656252857
2011 4.00472650244164
```

### -- correlation

To calculate correlation between star rating and helpful votes.

Query:

```
select corr(star_rating, helpful_votes) as correlation
from amazon_reviews_include
```

Output:

```
hive> select corr(star_rating, helpful_votes) as correlation
> from amazon_reviews_include
> ;
Query ID = hadoop_20200413094846_30f4a14c-2316-46e2-ad4e-88c9e5a224b8
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1586766680493_0011)
```

VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1 .....	container	SUCCEEDED	26	26	0	0	0	0
Map 5 .....	container	SUCCEEDED	26	26	0	0	0	0
Map 6 .....	container	SUCCEEDED	20	20	0	0	0	0
Reducer 2 .....	container	SUCCEEDED	26	26	0	0	0	0
Reducer 3 .....	container	SUCCEEDED	22	22	0	0	0	0
Reducer 4 .....	container	SUCCEEDED	1	1	0	0	0	0
Reducer 7 .....	container	SUCCEEDED	1	1	0	0	0	0

```
VERTICES: 07/07 [=====] 100% ELAPSED TIME: 200.75 s
OK
-0.022441850004671932
Time taken: 201.485 seconds, Fetched: 1 row(s)
hive> █
```

## CONCLUSION

We were able to make use of big data technologies to analyse the amazon reviews dataset, answer data exploratory questions, compare product categories and observe trends in metrics over time. We were successfully able to perform exploratory data analysis on AWS Athena and perform queries in Hive over the amazon reviews external table in HDFS. We used numerous aggregation functions over numerical columns like star rating and total votes to help us group it with other metrics and observe different trends according to the product categories. When we were trying to compare the average star ratings for product categories 'Video Games' and 'Digital Video Games', the average dropped drastically for digital video games around 2008 and then resurged but was still less than the almost consistent average star ratings for video games. This just demonstrates one of our findings. To conclude, we were able to perform detailed analysis on the amazon review dataset and we made use of hive queries to obtain results.

## REFERENCES

<https://piazza.com/class>

<https://aws.amazon.com/emr/>

<https://aws.amazon.com/emr/features/hive/>

[http://www1.udel.edu/evelyn/SQL-Class3/SQL3\\_Stat.html](http://www1.udel.edu/evelyn/SQL-Class3/SQL3_Stat.html)